

Data Structures: Part 1

Immutable state

- Major difference between Clojure and all other dialects of Lisp

Immutable state

What is it?

- Data is not modified
- So what happens when state is updated?
 - New data is created.
 - Old data still remains.

Immutable state

Why?

- Makes program easier to reason about (static analysis)
- Reduces dependence between functions: more modularity + code reuse
- Shorter and more legible code == less bugs.
- Everything is threadsafe, allowing for implicit parallelism (no need for locks)

Immutable state



- Alan Kay, Lisp hacker and inventor of Object-Oriented Programming

Immutable state

Does it affect performance?

- The short answer: no!
- Garbage collection on the JVM is now highly optimized
- Memory usage is controlled by eliminating intermediate copies of data
 - In other languages as a compiler optimization (fusion)
 - In Clojure by not creating them in the first place (transducers...more on this later)

Immutable State

In Clojure immutability is encouraged, but not enforced.

- Functional programming avoids variable assignment
- Functions are referred to as "pure" or having "referential transparency" if:
 - No global state passed in
 - No local state
 - They cannot modify other parts of the program (referred to as "side effects")

- Pure functions are "idempotent"— applying multiple times to same input does nothing (like mathematical functions)
- Most built-in functions in Clojure are pure
- Impure functions end with a ! ("bang")

The Collection Abstraction

“ It is better to have 100 functions operate on one data structure than 10 functions on 10 data structures.”

- Alan Perlis, first recipient of the Turing Award

- Clojure extends the "list processing" paradigm to other data structures by abstracting over them with a consistent API
- All data structures implement the "Collection" abstraction ("coll" for short)

The Collection Abstraction

Abstracing data structures based on shared APIs allows Clojure to optimize them under the hood.

- Even the same data structure may implement functions differently based on the number of elements
- It's often faster to implement higher order functions with Java methods than recursive Clojure functions.

- The Clojure compiler is a mix of Clojure and Java code, with APIs managed by Java interfaces
- Clojure encapsulates this using parametric polymorphism in the form of *protocols* (covered later)
- It's important to understand what particular implementation is being used in order to write performant code, but more specific details are beyond the scope of this workshop
- If you're interested, there's a link to the source for every function in the cheatsheet. Don't be afraid, the Clojure compiler is very accessible!

Types of Collections

Lists

- S-expressions starting with `list` or `quote`
- Example: `'(1 2 3)`

Types of Collections

Vectors

- Ordered, like lists but more performant
- Square brackets
- Example: `["foo" "bar" "baz"]`

Types of Collections

Maps

- Associates keys with values
- Unordered, except for `sorted-map` and `array-map`
- Curly brackets
- Example:

```
{:language "Clojure", :creator "Rich Hickey"}
```

Types of Collections

Sets

- Unordered
- Elements can only appear once
- Curly brackets preceded with hashtag
- Example:

```
#{"clojure" "scheme" "haskell" "ocaml" "scala"}
```

- Not covered in this workshop

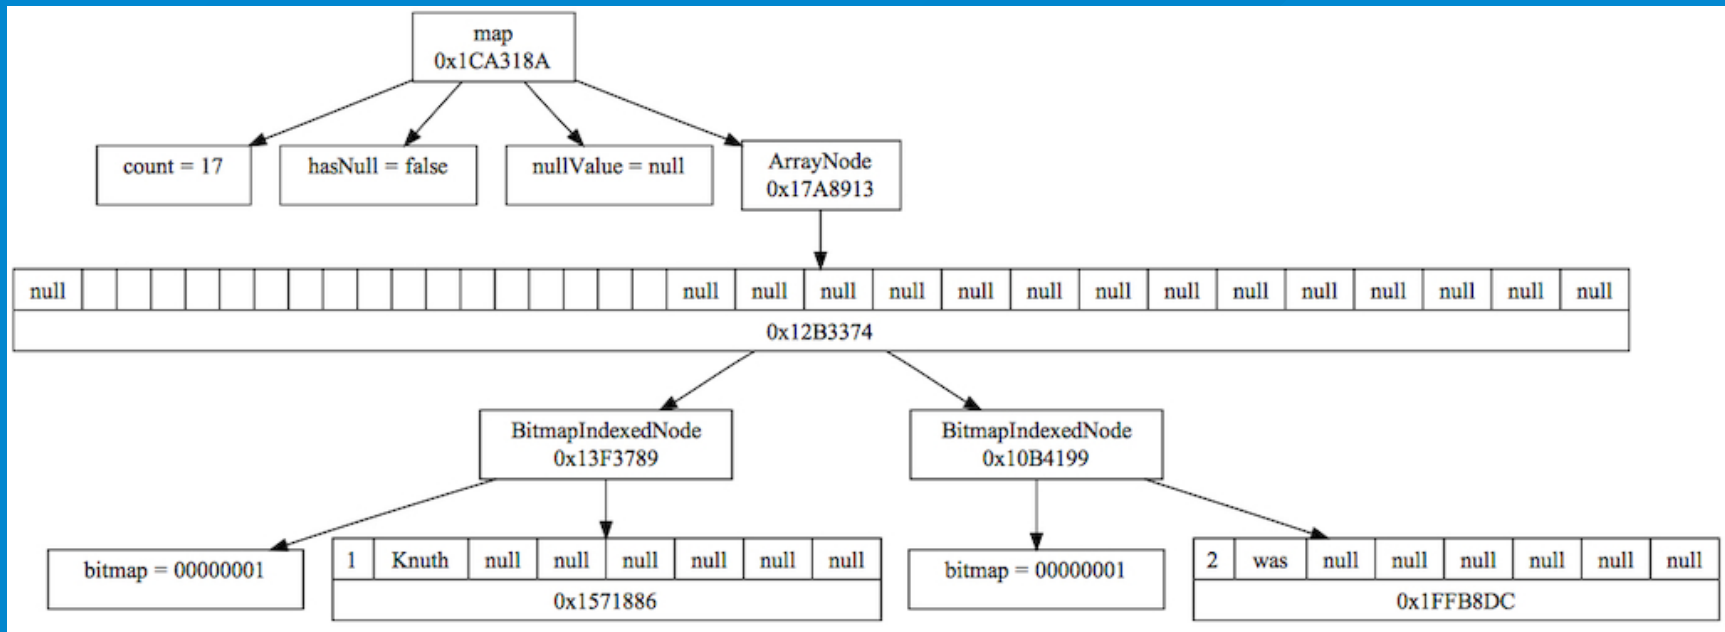
Under the Hood

All collections except for linked lists are "tries."

- Name comes from "retrieval"
- Officially pronounced like "tree," but commonly pronounced like "try" so as to avoid confusion
- Invented in 1959, but became popularized in Chris Okasaki's *Purely Functional Data Structures* (1998)

- Trees with data *only* stored in leaves
- Location in tree determines location (index or key) of value
- Grow dynamically when new values are added (preallocation unnecessary)

Under the Hood: Maps & Sets




Maps and Sets are "hash array mapped tries" (HAMTs)

HAMTs

- Invented by Phil Bagwell in "Ideal Hash Trees" (2000) then implemented in Scala
- Optimized for immutable key-value storage
- As name implies, keys are hashed to ensure constant length
- *Most of the time* faster for all common operations than hash-tables using buckets
- Lookup can often be done in constant time, although worst-case is linear

Under the Hood: Vectors



Closure vectors are 32-array binary search tries

Clojure Vectors

- Invented by Rich Hickey
- Binary structure means logarithmic time from root to leaf...
- ...except in this case each node has up to 32 children (unlike image above)...
- ...meaning \log_{32} time, which is *effectively* constant

- This makes large vectors extremely efficient, even into millions of elements
- Time complexity applies for most common operations: append, update, lookup, and taking subvectors
- However, *not* for prepending: unlike lists, new elements go at the end
- Vectors make ideal stacks: `conj`, `pop`, `peek`

Collection Functions

All collections share a basic API

- `=`: checks value equality of a collection compared to other collections
- `count`: returns number of elements in a collection
 - Takes linear time so use sparingly!

- `conj`: adds an item to a collection in the most efficient way (short for "conjoin")
 - Note: "most efficient" means in a different place for each collection type. Be careful!
- `empty`: returns an empty collection of the same type as the argument
- `seq`: gets a sequence of a collection

How to Work With Lists, Vectors, and Maps Coming Up in Part II...